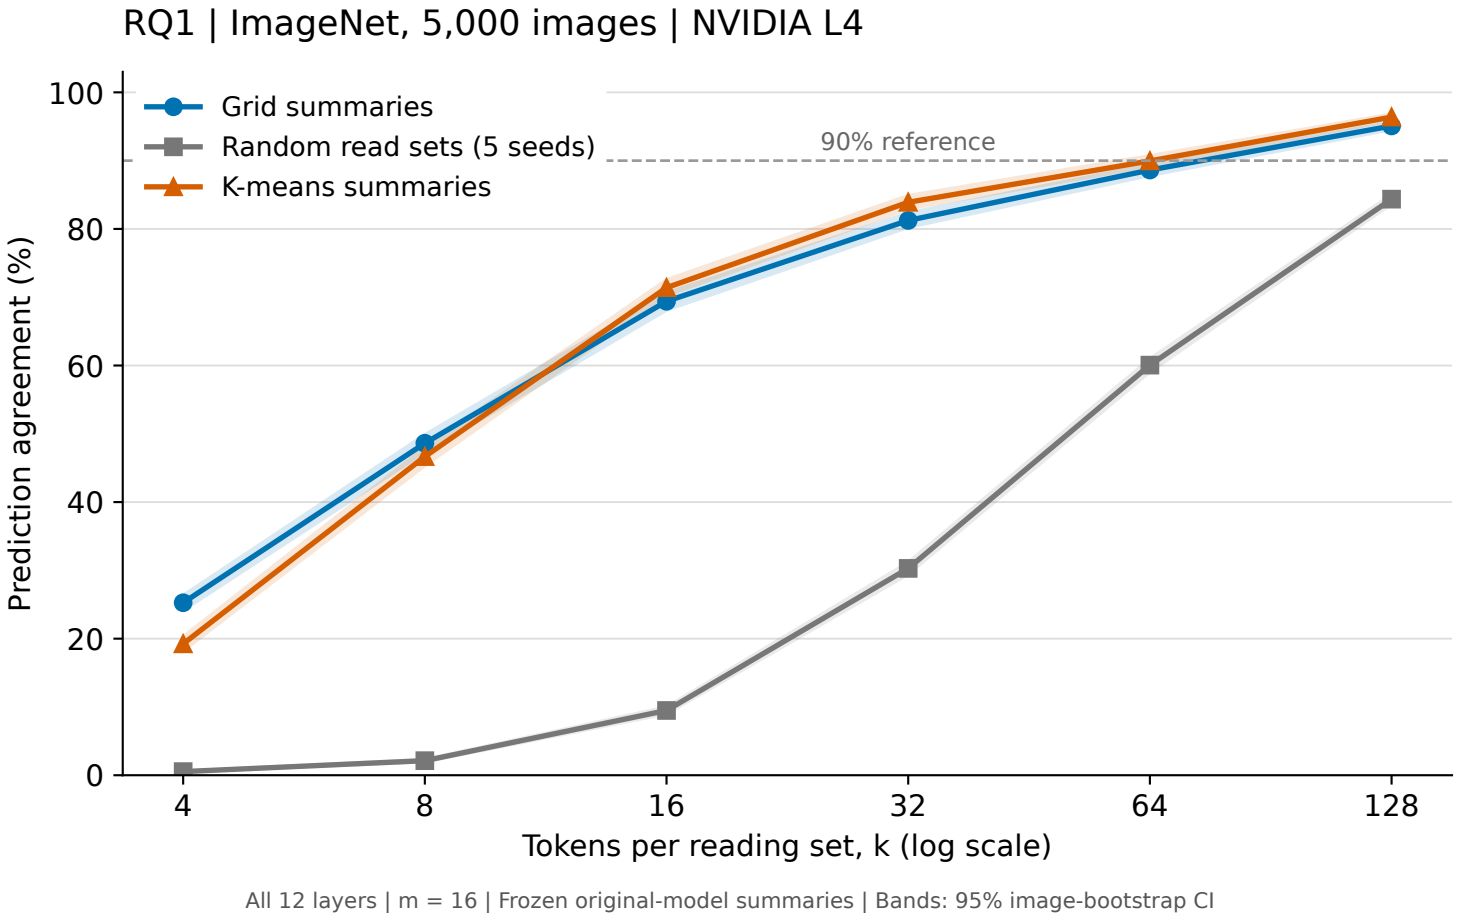
RQ1 | ImageNet, 5,000 images | NVIDIA L4
Grid summaries
Random read sets (5 seeds)
K-means summaries
90% reference
Prediction agreement (%)
100
80
60
40
20
0
4
8
16
32
64
128
Tokens per reading set, k (log scale)
All 12 layers | m = 16 | Frozen original-model summaries | Bands: 95% image-bootstrap CI