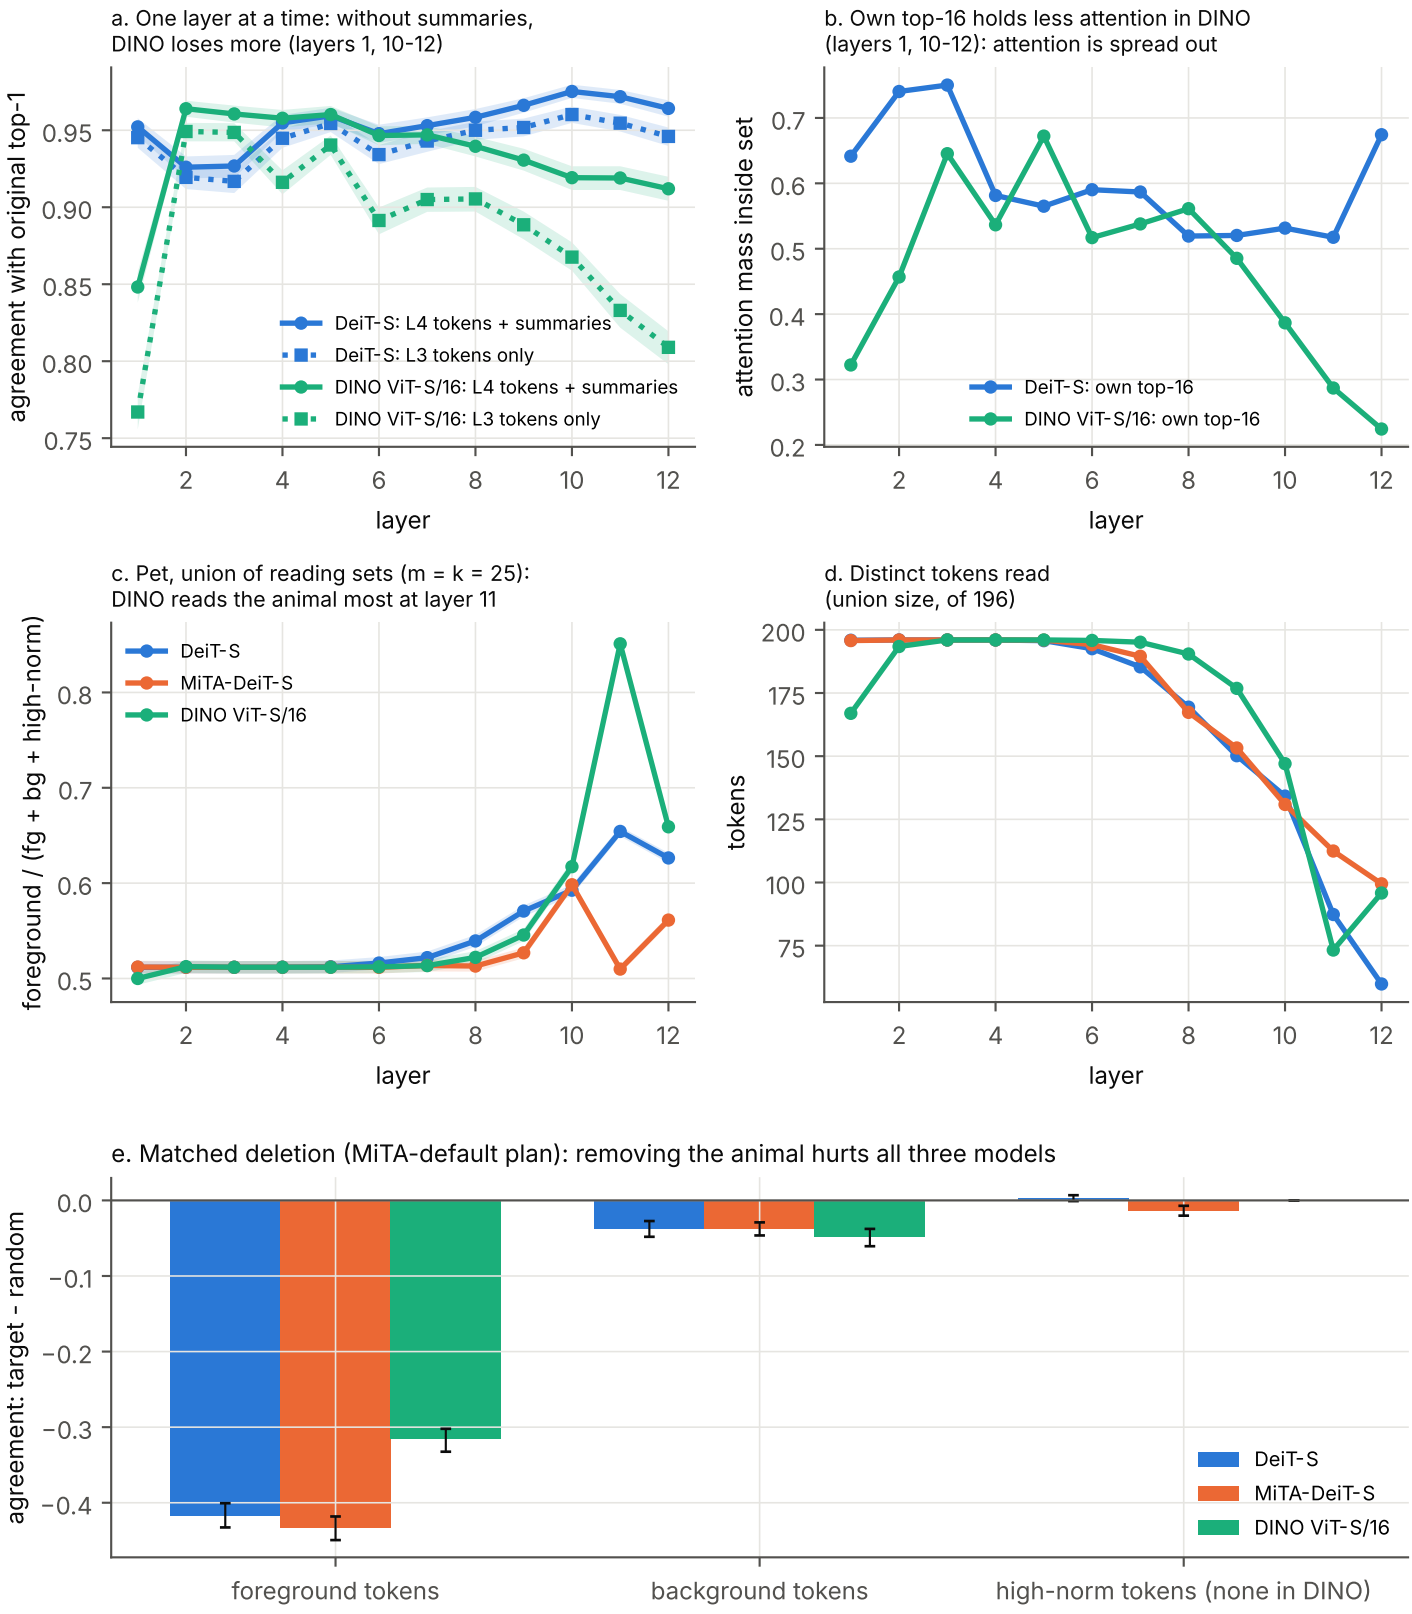
a. One layer at a time: without summaries,
DINO loses more (layers 1, 10-12)
agreement with original top-1
layer
DeiT-S: L4 tokens + summaries
DeiT-S: L3 tokens only
DINO ViT-S/16: L4 tokens + summaries
DINO ViT-S/16: L3 tokens only
b. Own top-16 holds less attention in DINO
(layers 1, 10-12): attention is spread out
attention mass inside set
layer
DeiT-S: own top-16
DINO ViT-S/16: own top-16
c. Pet, union of reading sets (m = k = 25):
DINO reads the animal most at layer 11
foreground / (fg + bg + high-norm)
layer
DeiT-S
MiTA-DeiT-S
DINO ViT-S/16
d. Distinct tokens read
(union size, of 196)
tokens
layer
e. Matched deletion (MiTA-default plan): removing the animal hurts all three models
agreement: target - random
foreground tokens
background tokens
high-norm tokens (none in DINO)
DeiT-S
MiTA-DeiT-S
DINO ViT-S/16