(a) A head as a per-image network

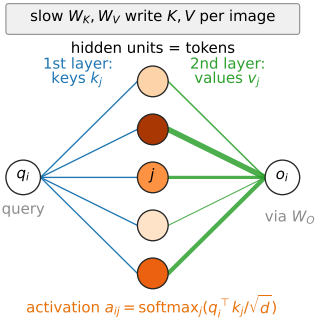

(b) Reading R0

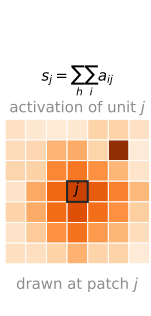

(c) Premises and corrections

premise	measured by	correction	stand-in
M1 activations are image-specific	mean-map share	C1 subtract mean map	subtract centre Gaussian
M2 unit j stands for patch j	position-probe accuracy	C2 relocate by probe	matched blur
M3 activation = contribution	R0 share on high-norm tokens	C3 weight by $\|W_O v_j\|$	zero high-norm tokens

each correction must beat its stand-in